

Basic Map Usage



Grant Little

<http://www.grantlittle.me>

grant@grantlittle.me

Overview

- What is a distributed map
- Adding data to a map
- Resiliency
- Persistent Storage
- Searching Maps and Predicates
- Indexing

Distributed Map

Node 1

Node 2

Node 3

Hazelcast API

Java Client

Ben

Brian

Claire

Jack

Joan

John

Julie

Mary

Michael

Simon


Data Distribution

Node 1



Partitions

Node 2



Partitions

Node 3



Partitions

Data Backups

Node 1

John (Primary)

Brian (Primary)

Julie (Backup)

Mary (Backup)

Node 2

Julie (Primary)

Joan (Primary)

John (Backup)

Simon (Backup)

Node 3

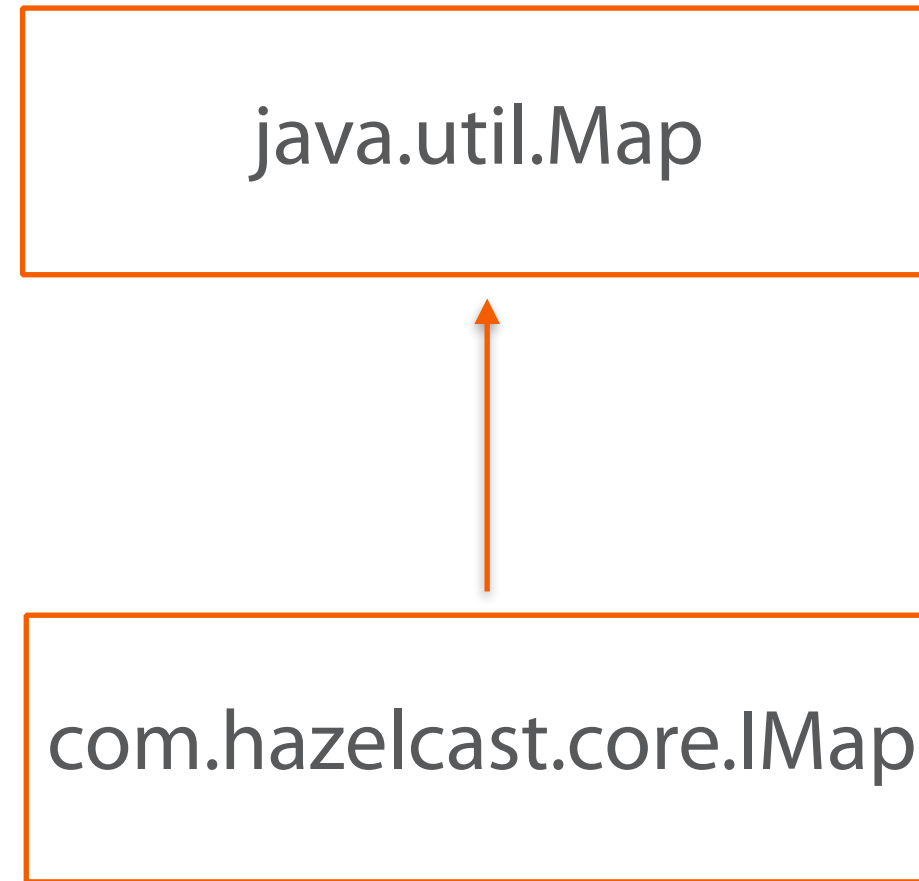
Simon (Primary)

Mary (Primary)

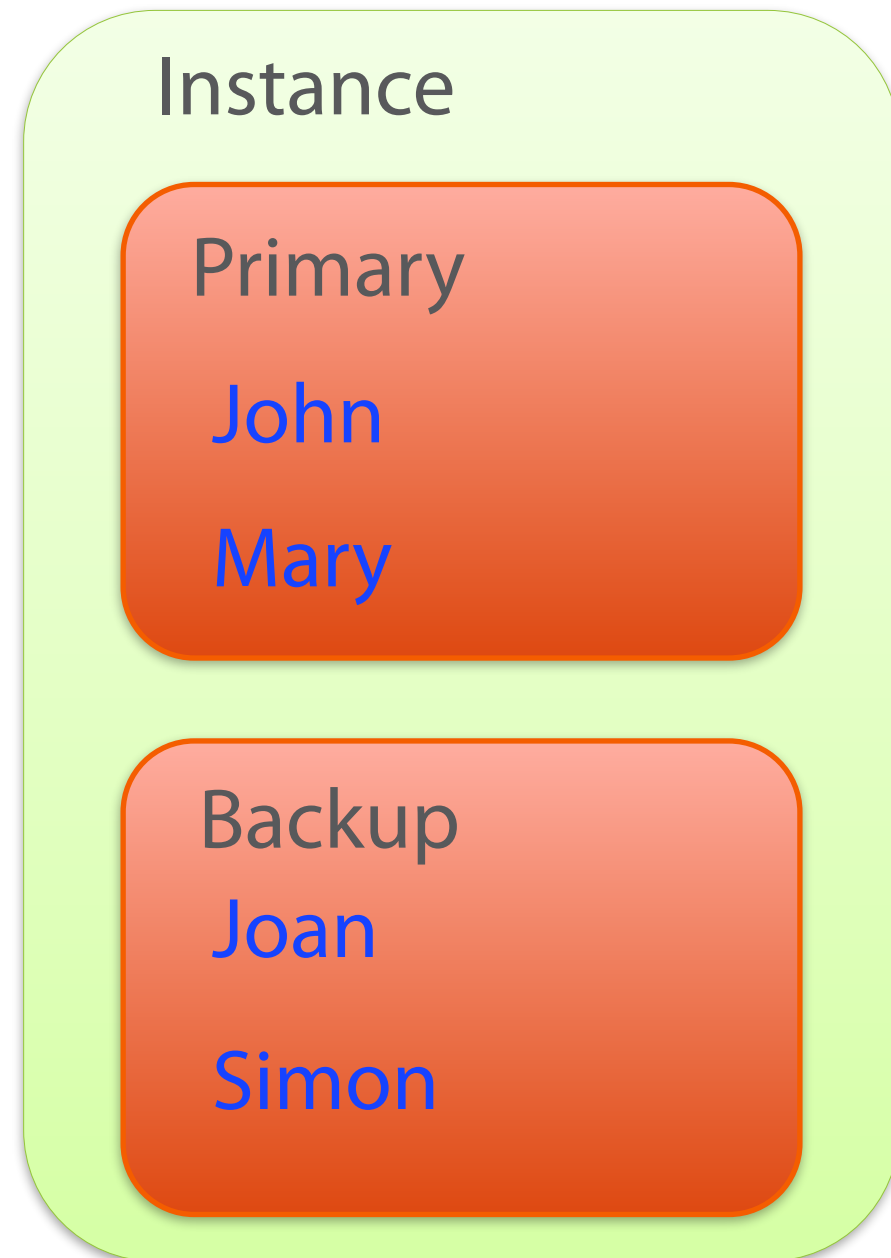
Brian (Backup)

Joan (Backup)

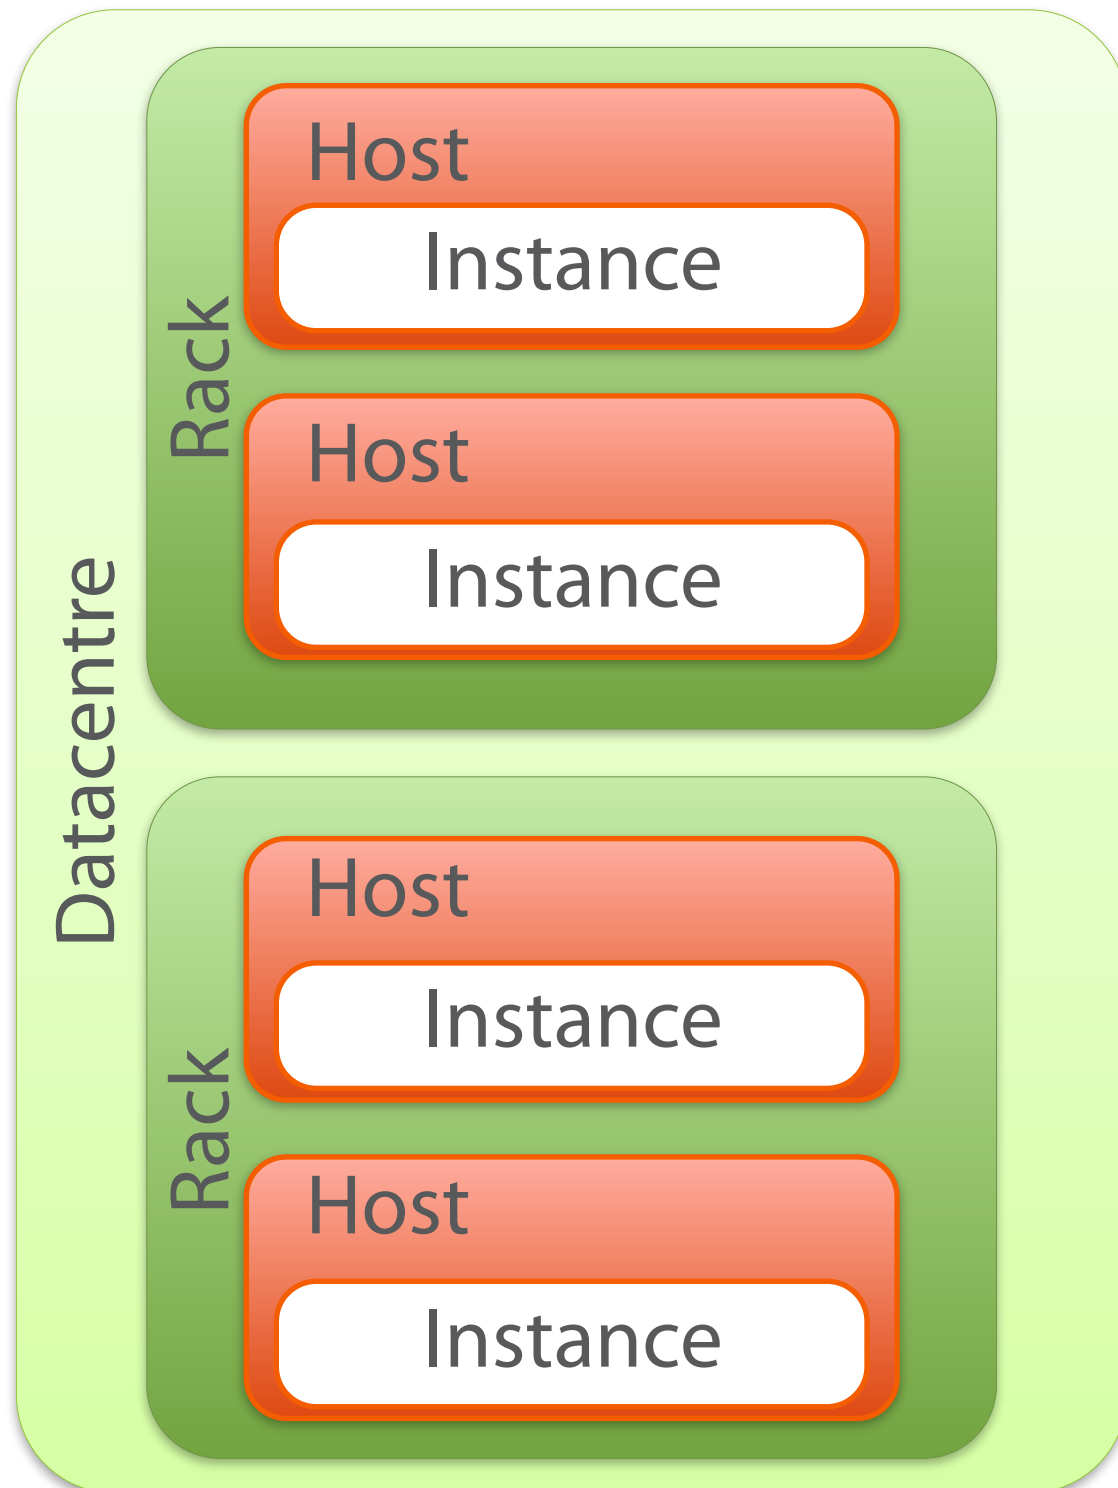
IMap Interface



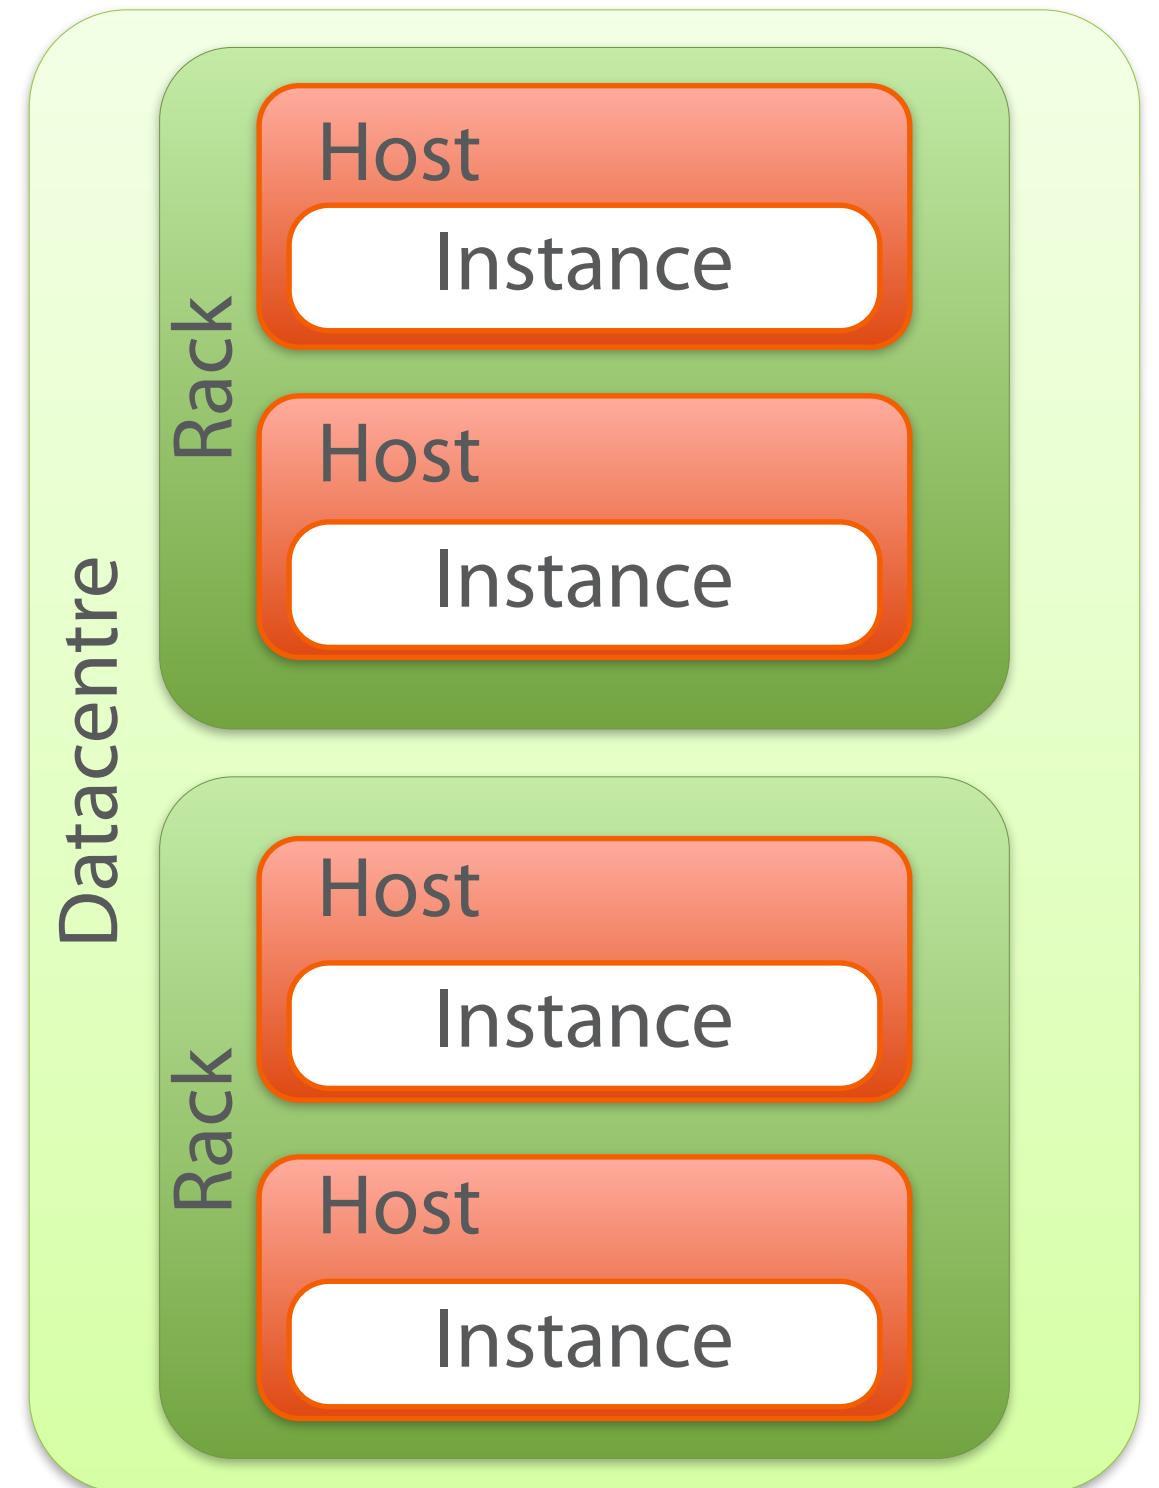
Activation of Backup Data



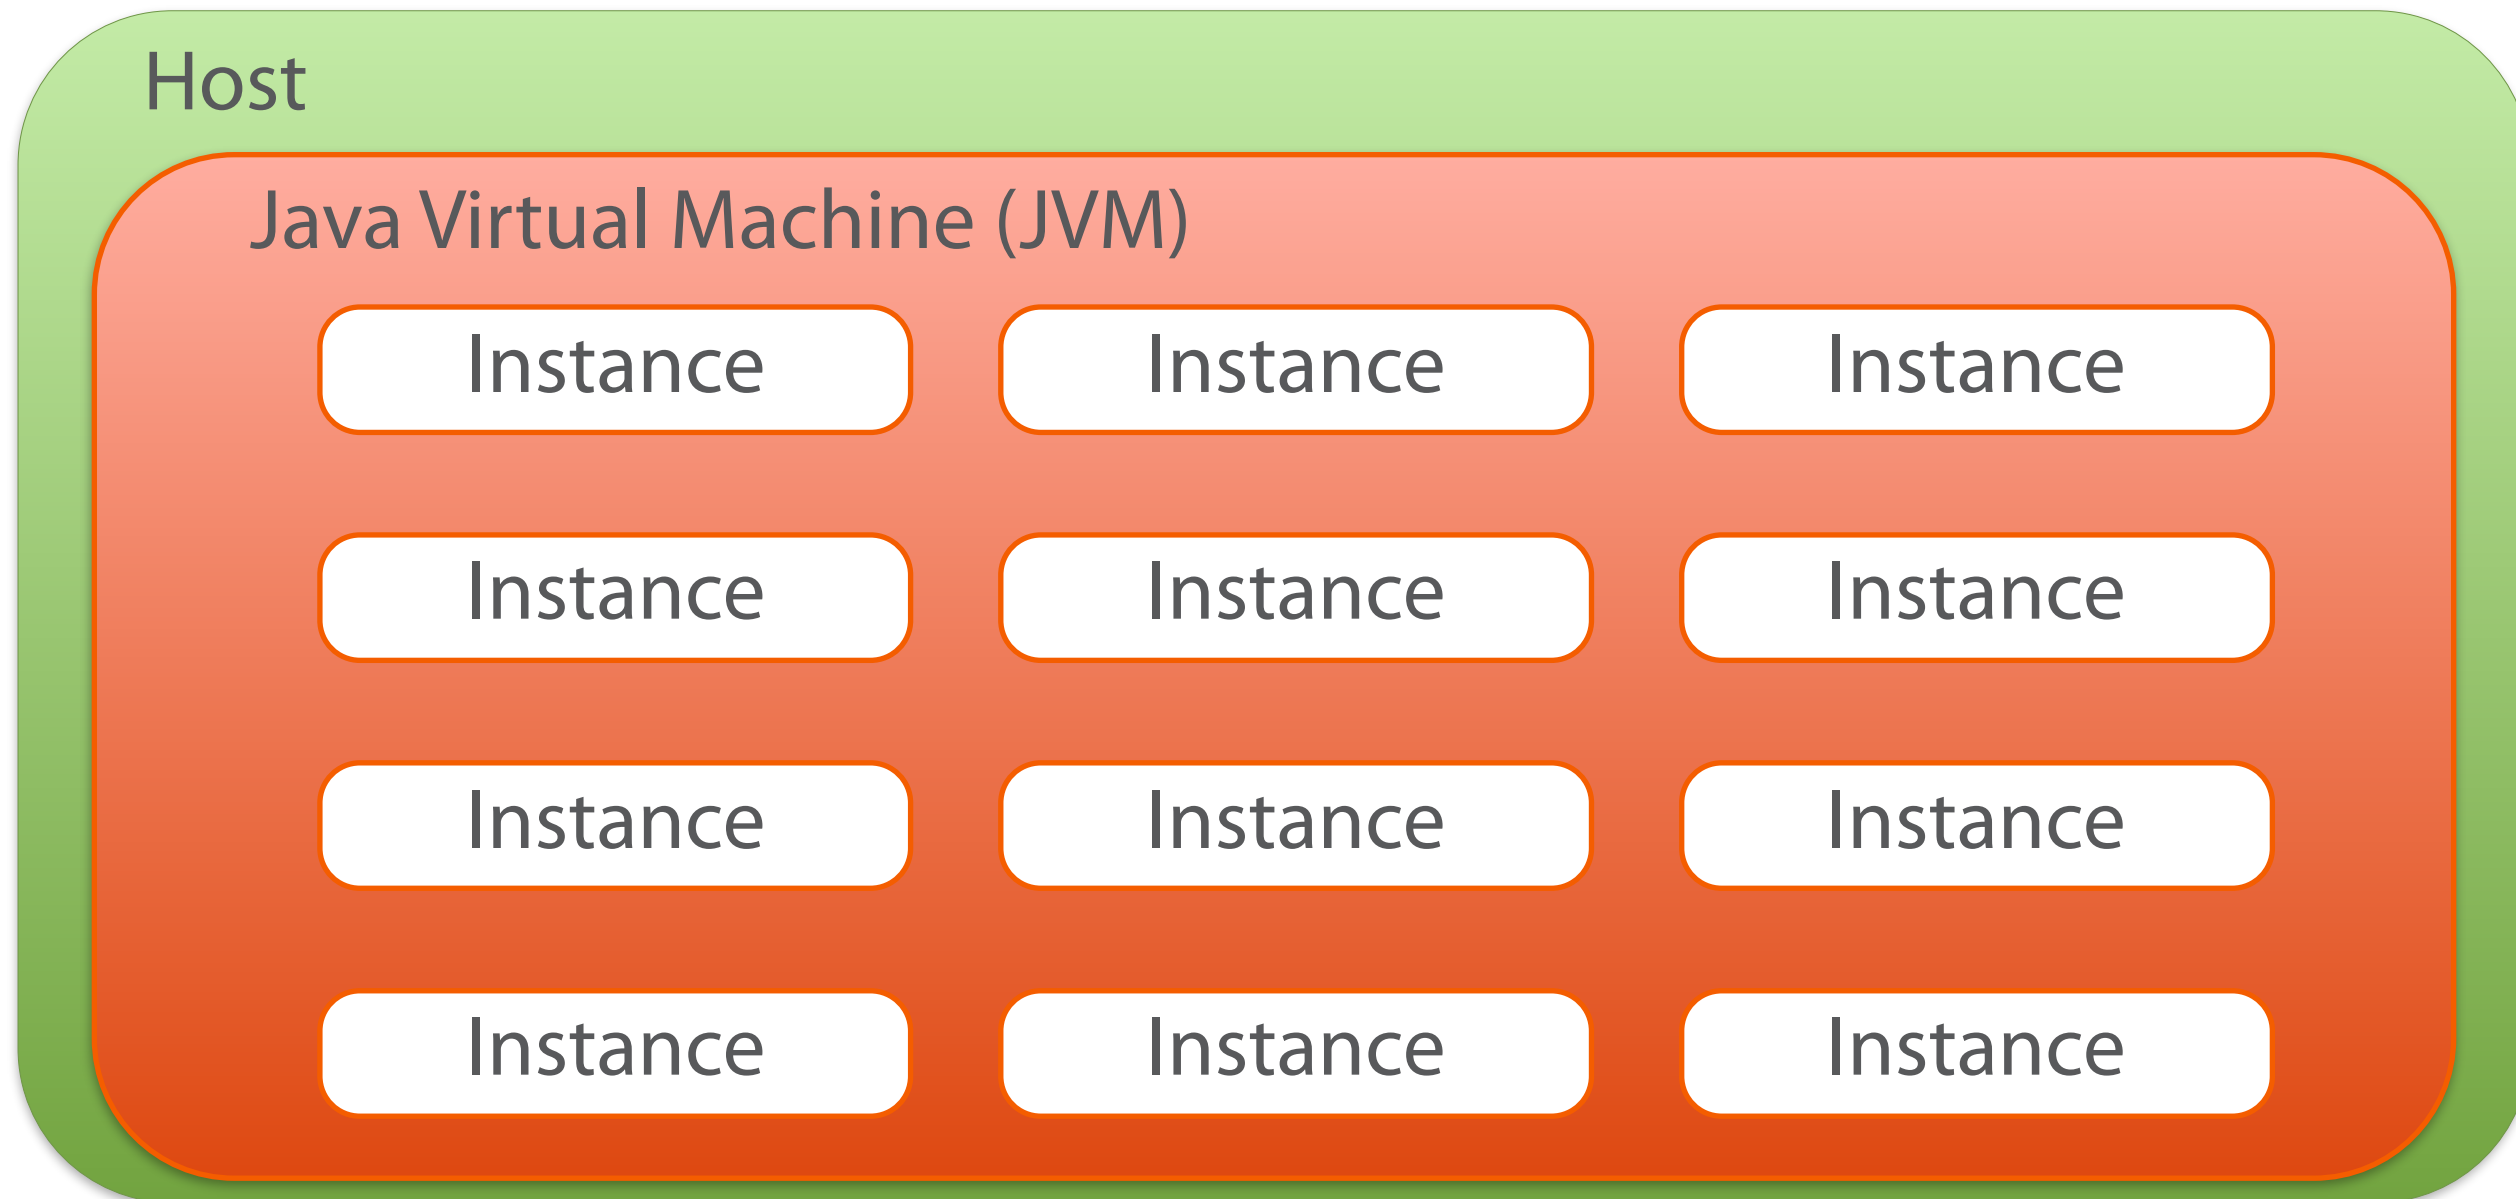
Possible Production Setup



1 Hazelcast Instance
=
1 Storage Node
=
1 Cluster Member



Possible Development Setup



1 Hazelcast Instance = 1 Storage Node = 1 Cluster Member

Map Store Persistence

- Maps store all data in memory
- When all storage nodes are shutdown the data is lost
- Use Map Stores for more persistent storage
 - Create a class that implements `com.hazelcast.core.MapStore`
 - Configure Hazelcast to use the `MapStore` for a particular map

Synchronous vs Asynchronous MapStore

- Synchronous
 - Every Map update results in the store or delete method on the MapStore being called
- Asynchronous
 - Every few seconds the Map Store is passed a batch of updates via storeAll or deleteAll methods on the MapStore.
 - Possible risk of data not being saved by the MapStore if the storage node is lost

Searching for Data

- Customer object has a “dob” (date of birth) attribute

```
private Date dob;
```

```
public Date getDob() { return dob; }
```

```
public void setDob(Date dob) { this.dob = dob; }
```

- Search for Customers who date of birth falls within a range

Main Predicates Available

- and
- between
- equal
- greaterEqual
- greaterThan
- ilike
- in
- lessEqual
- lessThan
- like
- not
- notEqual
- or
- regex

Or implement your own by implementing `com.hazelcast.query.Predicate`

SQL Predicates

Finding Data Using a SQL Like Language

SQL Predicates

- Possible to write queries in a SQL like language
 - `name = 'Grant'`
 - `name != 'Grant'`
 - `email NOT LIKE '%@pluralsight.com'`
 - `transactionAmount >= 10.00 and transactionAmount < 20.00`
 - `country IN ('US', 'AU', 'IRE', 'NZ', 'UK')`
 - `currency NOT IN ('AUD', 'EUR', 'GBP', 'USD')`
 - `transactionAmount BETWEEN (10.00, 20.00)`

Indexing

- Hazelcast stores map objects in their serialised form
- When searching:-
 - Each object is deserialised
 - The predicate is applied to the deserialised data
- Avoid the deserialisation by indexing

Review

- Distributed Map to store objects
- Insert, update and delete from a Map
- Hazelcast distributes data among the cluster members
- MapStore for persistent storage of map data
- Searching using Predicates and the SqlPredicate
- Indexes to improve search performance

Next “Working with Map Data”